PIA – Protein Inference Algorithms
Tutorial

Creation date 04.12.2017, 20:29:05.

Contents

1 Introduction

1.1 PIA – Protein Inference algorithms

PIA is an open source toolbox for MS based protein inference and identification analysis. As in bottom-up MS proteomics actually peptides are identified, but most often the entities of interest are proteins, the protein content of an analysed sample must be constructed from the knowledge of the contained peptides. This step is known as "protein inference" and is the heart of PIA. Furthermore, PIA can be used to inspect, analyse, perform quality checks on and filter identified peptide spectrum matches (PSMs) and peptides. This can be performed either via a web frontend, KNIME nodes or the command line. While the latter method is intended for advanced scripting and the command line can now also be downloaded as a Docker image, we will mainly discuss the usage of KNIME nodes in this tutorial.

1.2 Prerequisites

All data and workflows can be downloaded in the tutorial repository at
`https://github.com/julianu/pia-tutorial`.

- Knowledge of mass spectrometry based bottom-up peptide identification

- Basic knowledge of KNIME and constructing workflows with KNIME

- It is helpful to know the basics of OpenMS for KNIME, mainly the usage of spectrum identification nodes. For the advanced and quantification chapter, also the basics of quantification should be known. But both are also shortly sketched in this tutorial. (The openMS tutorial can be found at `http://www.openms.de/tutorials/`)

- The tutorial was tested using the stable PIA 1.3.5 nodes and the stable OpenMS 2.1.0 nodes with KNIME 3.4.2

1.3 Version

This tutorial was created on 04.12.2017 at 20:29:05.

2 Installation of PIA KNIME nodes

If not yet done, you first need to download and install KNIME to your system (`https://www.knime.com/downloads`). The easiest way to work with PIA inside KNIME is to install the **"KNIME Analytics Platform + all free extensions"**, which comes with many nice nodes and functions, including the PIA nodes.

If you did install KNIME without all free extensions, or upgraded from an older version, you can install the nodes from the community contributions repository. For this start KNIME and go to **HELP > INSTALL NEW SOFTWARE...**. Select the community contributions repository under the **WORK WITH** drop down menu, it should have an address similar to `http://update.knime.org/community-contributions/trusted/3.4`. The PIA nodes can be found in the **BIOINFORMATICS & NGS** group or simply by searching for them. Select the PIA nodes, click next, accept the license and restart KNIME after the installation is finished.

If all went well, you will see the PIA octopus on the splash screen of KNIME (together with all the other icons) and you will find the PIA nodes inside the **COMMUNITY NODES** in the Node Repository (usually left bottom side of screen). This tutorial also needs the OpenMS nodes (also in the community contributions' **BIOINFORMATICS & NGS**, next to PIA nodes) and the JavaScript Views nodes (regular KNIME updates repository, in the **KNIME LABS EXTENSIONS**) to be installed. If you installed **"KNIME Analytics Platform + all free extensions""** as recommended above, you need only to install the JavaScript nodes for the advanced tutorials.

3 Basic PIA analysis in KNIME

First download the workflows and data for the tutorial at `https://github.com/julianu/pia-tutorial`.

3.1 First workflow

First, we will run a minimal workflow identifying spectra in an mzML file with the search engine X!Tandem and using PIA for the protein inference and also analysis of the identified PSMs and peptides.

Import and open the workflow **01-PIA_FIRST_ANALYSIS** from the provided workflows into your KNIME workspace. Please select the file **1A-QEXHF04026.MZML** as the analysed mzML file (upper **INPUT FILE**) and the provided FASTA file with decoys (**UNIPROT-PROTEOME-MOUSE-SPIKED-CRAP-IRT-ISOS-2017_03-DECOY.FASTA**) as database for spectrum identification. We will use X!Tandem for spectrum identification in this workflow. If you like, you can change it to any other supported spectrum identification software in KNIME or make the identification externally (e.g. with Mascot or Proteome Discoverer. But then, you need to load the results directly into the **PIA COMPILER** node, see below.) The basic configuration settings for spectrum identification of the files used in the tutorial are:

- 10 ppm precursor tolerance

- 0.02 Da fragment tolerance (or 20 mmu)

- tryptic digestion with up to 2 missed cleavages

- fixed Carbamidomethyl of C and variable Oxidation of M of modifications

The **PEPTIDEINDEXER** is needed to add the protein / accession information to the identified peptide spectrum matches (PSMs). For the tutorial data, use "DECOY_" as the decoy string and prefix as decoy position and make sure that "Trypsin" is set as enzyme with "none" as specificity. After adjusting the **INPUT FILE** nodes (and maybe checking the other nodes' settings), run the workflow until the **PIA COMPILER**.

This node takes a list of files. Note, that the ports from OpenMS need to be converted into a URI file list. You could alternatively use a **LIST FILES** node to select the files containing prior performed spectrum identifications in any supported format like Mascot Dat, X!Tandem's XML files, Proteome Discoverer's MSF files or any mzIdentML file, and use these files as input for the compiler. The **PIA COMPILER** is necessary to structure

the data before a PIA analysis. It must be performed only once per set of identification files. The node has as configuration settings only a compilation name, which can be chosen freely. The number of files passed to PIA is not limited, though processing more files needs more main memory. Each file in the compilation gets a FileID, which you can explore with some additional information of the compilation when looking at the **VIEW: SUMMARY** after running the **PIA COMPILER**. These IDs can also be used for filtering and advanced settings later on.
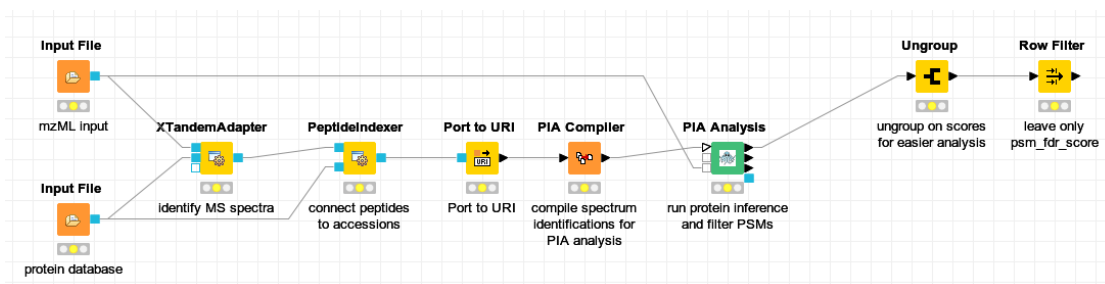
Figure 1: A basic protein analysis workflow using X!Tandem for spectrum identification and PIA for protein inference.

The first connected port of the **PIA ANALYSIS** node is for the PIA compilation, directly coming from the compiler node. If you should have a saved compilation, you could also load this with an **INPUT FILE** node together with the second input port. The first port of both with a suitable configured file will be processed for analysis. The third input port is used to pass spectrum data to PIA, which can be used to visualise automatically annotated spectra. If you don't want to use the spectrum viewer, you should not connect anything to this port. As the matching of the PSMs to their spectra might take some time, you should consider using this only when you need a spectrum annotation.

3.1.1 General Settings

After running the compiler, open the settings of the **PIA ANALYSIS** node, either by double clicking on the node or right clicking and selecting **CONFIGURE...**. You will find four tabs for the settings: one general and one for each of the three levels of analysis (i.e. PSMs, peptides and proteins). If you connected the compiler node directly to the analysis node, select the column containing the PIA XML file in the respective setting (there is mostly only one named "gzipped PIA XML file"). You can also set, whether PIA should fail if no decoys were found in the analysis, whether PSM sets should be created and whether modifications should be considered to distinguish peptides (see Figure 2). Creating PSM sets should always be used, if the same spectrum file was analysed with different search engines. This

option will then combine the results of multiple searches. Otherwise, when e.g. combining the results of multiple MS runs identified by only one search engine or, like in the example workflow, having only one search engine for one file, it can savely be deselected. Mostly, a peptide and it's scoring influence for the proteins should only be described by its sequence. If you have stable modifications, though, you can select to distinguish peptides by sequences and modifications. The **PIA ANALYSIS** node allows to export the analysis into several file formats. The level of the export (PSM, peptides and protein) as well as the format can be selected appropriately. Currently only one export file can be created at once.

Figure 2: The general settings dialog of the **PIA ANALYSIS** node.

The analyses at the PSM and peptide level of PIA will be performed only for the input file given by the FileID in the respective settings tab. Usually, the IDs start with 1 and are sorted by the order of the given input files. A special case is the combination of all runs (either with PSM sets or without), also called overview, which has always the ID 0. The 0 for FileID is the default for the PSM and peptide level analyses.

3.1.2 PSM Settings

Now, have a look at the PSMs settings of the **PIA ANALYSIS**. The first setting is the just mentioned file ID, which is 1 in the example and thus reflects the first file of the compilation (Note: we have only one file in this example, but still 0 could be used for the overview, if e.g. you would like to calculate the "Combined FDRScore". Then you would also need to enable the creation of PSM sets and the calculation of the "Combined FDRScore". For one single input file, though, this would be identical to the normal FDRScore). Remember, you can check which file has which ID in the view of the compiler node. Next you can choose to calculate the false discovery rate (FDR), and thus also FDRScore and q-value, for all input files. The FDRScore [2] smoothes the FDR q-values in an analysis and thus facilitates a better discrimination of identifications instead of using the FDR q-values alone. If PSM sets

are created, also the Combined FDRScore can be calculated, which furthermore allows the combination of search results from multiple MS runs as well as identifications from different search engines.

Figure 3: The PSMs settings dialog of the **PIA ANALYSIS** node.

A very important step is the selection of how decoys are distinguished from target identifications. PIA allows to use regular expressions for this, which are applied to the accessions. In the example in Figure 3, "DECOY_.*" is set as regular expression. So each accession starting with the string "DECOY_" (and all its peptides) will be assigned to be decoys and all not matching accessions to be targets. Alternatively, if "by search engine" is selected, identifications must be annotated in the input files as targets and decoys. This is e.g. the case if in Mascot the "Decoy" option is selected for an MS/MS search. For better compatibility though, the usage of a target-decoy database and the assignment by regular expressions is recommended. For this to work, don not filter the identifications for decoys or FDR before passing them to PIA.

Some search engines report more than one identification per spectrum. In the next option, you can choose to either use all these identifications or only the one with the best score for

the FDR analysis and all following steps. Finally, you can select which scores are used for FDR estimation. If a search engine reports multiple score (e.g. X!Tandem's Hyperscore and Expect), you can choose the preferred score here. If no score was chosen, PIA will use the main score of the search engine (or, if this was not given, just any score, but this consistanly).

All the settings up to this point are used for the FDR calculation (except for the output file ID) and influence the peptide and protein reports as well. The filters though afflict only the PSM level and its report. Here you can chose from the available filters and set the parameters accordingly. For score filters it is necessary to select the according score as well in the line below the actual filters selection. After selecting a filter and setting the parameters, don't forget to click **ADD**. The currently activated filters will be shown in the list. In the example workflow, a filter for the "PSM FDRScore $<= 0.01$" is set. Keep in mind, that all filters have to be fulfilled for a PSM to pass the filters (the filters are additive).

3.1.3 Peptides Settings

Next, select the peptides settings. All settings here are only relevant for the peptide export and do not afflict the protein inference. Therefore, you can also turn the peptide inference off and will get an empty reporton the peptide level. This might be a good idea, if you want to save time and main memory during the analysis. The selection of the file ID has the same meaning as on the PSM level: exporting only information of the given file, or 0 for the combination of all results. Again: if in doubt, check the file ID at the compiler view. Also the filters are used in the same way. An exception are the PSM level filters: with these, the PSMs which are actually used to create peptides, are filtered before the epptide inference. In the example workflow and in Figure 4, only PSMs with "PSM FDRScore $<= 0.01$" are inferred to peptides, all other PSMs are discarded.

Figure 4: The peptides settings dialog of the **PIA ANALYSIS** node.

3.1.4 Protein Inference and Protein Settings

Finally, take a look at the proteins settings. Also the protein inference can be turned off, if you are only interested in analyses on the PSM or peptide level. First, you have to chose an inference method. PIA provides you with three different methods (for a more thorough explanation of these methods, please refer to [1])

- **Occam's Razor** is based on the parsimony principle and returns the smallest set of proteins, which explain all identified peptides. This is a very widely used strategy for protein inference.

- **Spectrum Extractor** is the recommended method. It is also based on the parsimonious approach, but assigns a spectrum to only one peptide. These peptides are selected in such a way, that they increase the score of the most probable, possible protein group.

- **Report All** This strategy reports all possible protein groups, based on the given data. Use this with caution and only when you know, what you are doing! It's mainly used when searching for a special protein group in your data.

After selecting the inference method, you can apply a variety of filters, which will be applied on PSM, peptide and protein level and directly afflict the results of the inference. You should almost always filter on the PSM level FDR, either using the FDR Score or the Combined FDR Score, usually on a value of about 0.01. Remember: if you searched with only one search engine, the FDR Score will be the score of choice, for multiple search engines the Combined FDR Score. But you could also set filters which allows only protein groups with at least two peptides, and many more filters.

Next, you need to select the scoring method. Here, you should use "additive scoring" only if your base score has a "higher score is better" probability, like e.g. the Mascot Ion Score or X!Tandem's Hyperscore. Otherwise, use one of the other two scorings. The "multiplicative scoring" takes the number of identified peptides into account, in the way that the final protein score is usually better with more distinct peptides. The "geometric mean scoring" on the other hand calculates the mean of all peptide scores.

Finally, you need to select the base score and whether only the best PSM (recommended) or all PSMs of a peptide should be used for scoring. In our example, the FDR Score is used. Note: As base score you can only select scores on the PSM level, as the peptides will be generated during the inference according to your applied filters.

The protein report can also be filtered in the same way as the PSMs and peptides report. Be aware, that this is significantly different to setting an inference filter on protein level:

Figure 5: The proteins settings dialog of the **PIA ANALYSIS** node.

filters applied for the inference make it possible to not even create a protein group. Filters on the protein report afflict only what will be reported, not what will be inferred.

3.1.5 A word on filters

Though PIA provides you with many filters on each of the PSM, peptide and protein level, you can also apply these filters later in the workflow. For this, you can simply apply an appropriate **ROW FILTER** node after running the **PIA ANALYSIS**. But also keep in mind that setting the inference filters on the protein level are behaving very differently, as explained in the prior paragraph.
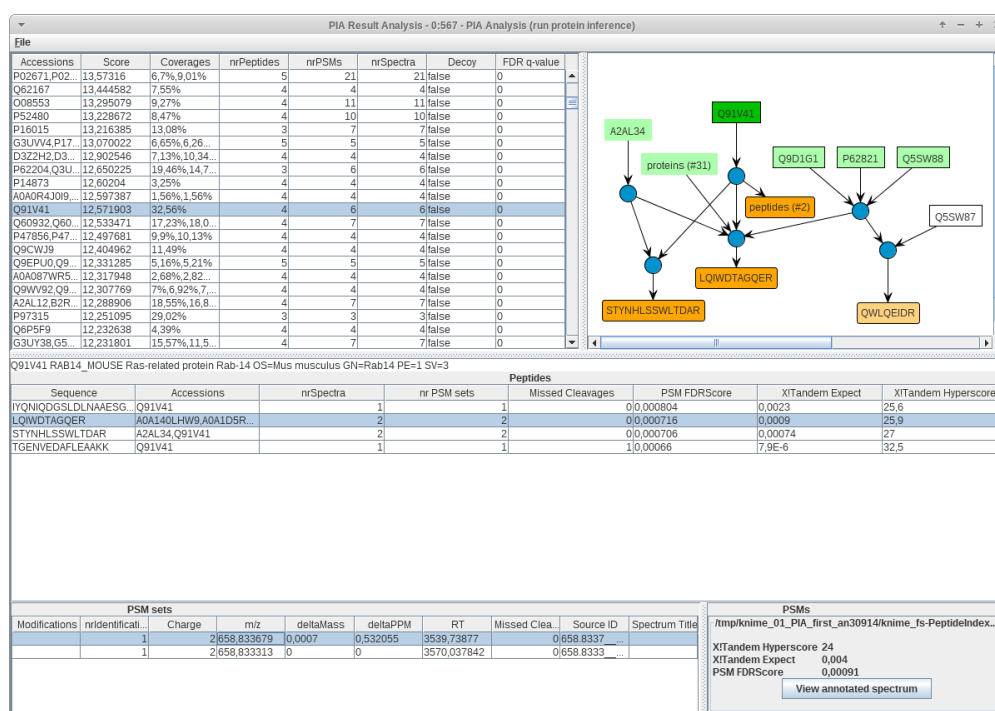
3.2 Looking at the Results

Now run the **PIA ANALYSIS** node. This will, depending on the loaded data, take some time. The example data should be processed in a few minutes though. The node has four output ports: the first two are the (filtered) reports on the PSM and peptide level for the selected input file, the third is the protein level report and the last is a file port for the exported data, if any was created according to the settings in the general settings tab.

3.2.1 Analysis Viewer

To explore the results of the analysis, right click on the **PIA ANALYSIS** node after the execution finished and select the **VIEW: PIA RESULTS ANALYSIS**. In the top left corner you will see all the inferred protein groups. PIA always works with protein groups, even if such a group might contain only one accession. All the accessions in one group have exactly the same evidence, i.e. the same PSMs and peptides, and cannot be distinguished on the given data and applied inference settings, at least without any further knowledge. The score is calculated using the selected base-score and scoring method. A higher score is always better (for base scores with "lower score better" a -log value is used for transformations). If the complete protein sequences were provided, the coverages for the proteins are calculated. Furthermore, the number of assigned spectra, PSMs and peptides are listed next to these. If the FDR was calculated, also the decoy status and the protein level FDR q-value is given.

For the currently selected protein group, the assigned peptides (i.e. the ones, that were not filtered out by the inference filters) are listed together with their information in the central list of the window. On the bottom left, all PSM sets of the currently selected peptide are listed and on the bottom right finally the individual PSMs are given. If the corresponding spectrum file was passed to the analysis node, you can view the annotated spectrum when clicking on the button by the PSM (see also PSM Results).

On the top right you will see a directed graph showing the relations between accessions of the currently selected protein group and its peptides and PSMs. This might look a bit overwhelming at first, but can be helpful, if you know how to read it. In the example workflow, take a look at the group "P02675" (it has a score of 46.66) and "Q91V41" (score of 12.57). The accessions of the selected group are coloured in dark green with a black border. Accessions of (not reported) sub-groups are given in light-green without border and accessions reported in other groups in light green with black border. Peptides are coloured in orange (dark for the selected group, light for other groups in the same way as accessions). The blue circles are drawn only to construct a correct tree. Nodes which hold multiple items, can be expanded by double clicking on them, as can peptides to show their corresponding

PIA Result Analysis - 0:567 - PIA Analysis (run protein inference)

File

Accessions	Score	Coverages	nrPeptides	nrPSMs	nrSpectra	Decoy	FDR q-value
P02671,P02...	13,57316	6,7%,9,01%	5	21	21	false	0
Q62167	13,444582	7,55%	4	4	4	false	0
O08553	13,295079	9,27%	4	11	11	false	0
P52480	13,228672	8,47%	4	10	10	false	0
P16015	13,216385	13,08%	3	7	7	false	0
G3UVV4,P17...	13,070022	6,65%,6,26...	5	5	5	false	0
D3Z2H2,D3...	12,902546	7,13%,10,34...	4	4	4	false	0
P62204,Q3U...	12,650225	19,46%,14,7...	3	6	6	false	0
P14873	12,60204	3,25%	4	4	4	false	0
A0A0R4J0I9...	12,597387	1,56%,1,56%	4	4	4	false	0
Q91V41	12,571903	32,56%	4	6	6	false	0
Q60932,Q60...	12,533471	17,23%,18,0...	4	7	7	false	0
P47856,P47...	12,497681	9,9%,10,13%	4	4	4	false	0
Q9CWJ9	12,404962	11,49%	4	4	4	false	0
Q9EPU0,Q9...	12,331285	5,16%,5,21%	5	5	5	false	0
A0A087WR5...	12,317948	2,68%,2,82...	4	4	4	false	0
Q9WV92,Q9...	12,307769	7%,6,92%,7...	4	4	4	false	0
A2AL12,B2R...	12,288906	18,55%,16,8...	4	7	7	false	0
P97315	12,251095	29,02%	3	3	3	false	0
Q6P5F9	12,232638	4,39%	4	4	4	false	0
G3UY38,G5...	12,231801	15,57%,11,5...	4	7	7	false	0

Graph labels: A2AL34, Q91V41, proteins (#31), Q9D1G1, P62821, Q5SW88, peptides (#2), Q5SW87, LQIWDTAGQER, STYNHLSSWLTDAR, QWLQEIDR

Q91V41 RAB14_MOUSE Ras-related protein Rab-14 OS=Mus musculus GN=Rab14 PE=1 SV=3

Peptides

Sequence	Accessions	nrSpectra	nr PSM sets	Missed Cleavages	PSM FDRScore	X!Tandem Expect	X!Tandem Hyperscore
IYQNIQDGSLDLNAAESG...	Q91V41	1	1	0	0,000804	0,0023	25,6
LQIWDTAGQER	A0A140LHW9,A0A1D5R...	2	2	0	0,000716	0,0009	25,9
STYNHLSSWLTDAR	A2AL34,Q91V41	2	2	0	0,000706	0,00074	27
TGENVEDAFLEAAKK	Q91V41	1	1	1	0,00066	7,9E-6	32,5

PSM sets

Modifications	nrIdentificati...	Charge	m/z	deltaMass	deltaPPM	RT	Missed Clea...	Source ID	Spectrum Title
	1	2	658,833679	0,0007	0,532055	3539,73877	0	658.8337...	
	1	2	658,833313	0	0	3570,037842	0	658.8333...	

PSMs

/tmp/knime_01_PIA_first_an30914/knime_fs-PeptideIndex...

X!Tandem Hyperscore 24
X!Tandem Expect 0,004
PSM FDRScore 0,00091

View annotated spectrum

Figure 6: The PIA analysis viewer allows an intuitive exploration of the data.

PSMs. All items can be re-arranged by drag-and-drop.

3.2.2 PSM Results

To look at the PSM level report-table, right click on the **PIA ANALYSIS** node and select **0: PSM RESULTS**. In this table you have almost all available information for the PSMs, like the amino acid sequence, a list of accessions, modifications, precursor charge, m/z value, the mass error, etc. Also, you have three columns for the scores (scores, score names and score shorts), which are lists each. Here, the score on a given position in the list corresponds to the score name and its short (which is an abbreviated name) on the same positions in the list. This makes it possible to export multiple scores in one row. In our example, we have three different scores per PSM: the X!Tandem Hyperscore and Expect, as well as the FDR Score.

For easier analysis, the scores can be ungrouped using the **UNGROUP** node on the PSM report, which is the second last step in the workflow. Take a quick look at the **UNGROUP**'s settings and verify, that it is set to ungroup on the three score columns (and not the accessions). Run the node and look at the results: the score columns contain now individual names and numbers, which can be used for further sorting and filtering. Note also,

that there is three times the number of rows after the ungrouping. To leave only the PSM FDR Score in the table, the final **ROW FILTER** is used.

An alternative, though not that comprehensive way to look at all reported PSMs (but not any PSM sets) is using the "PSM Spectrum Viewer". Here you can see the PSMs together with the automatically annotated spectra (using [3]). This only works, if you connected the identified spectrum file to the third PIA input port.
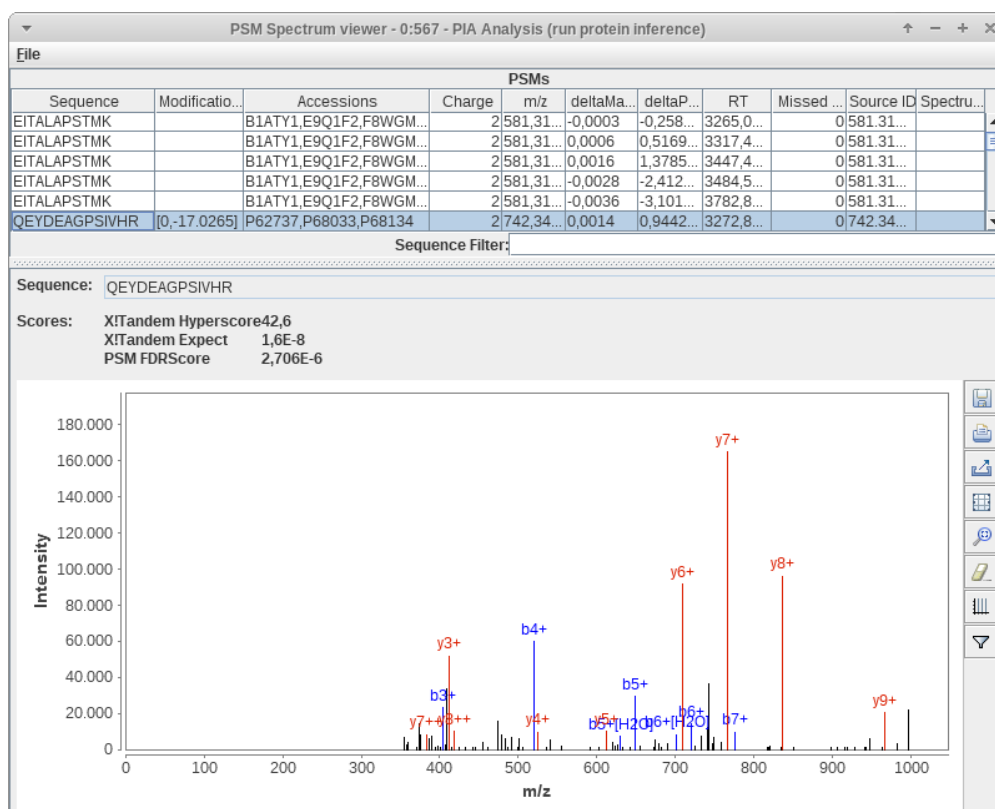
Figure 7: The PSM spectrum viewer showing an automatically annotated spectrum.

3.2.3 Peptide Results

The peptide results are given in the same way as the PSM results. The peptides are inferred either with or without taking modifications into account, as explained before.

3.2.4 Protein Results

The created protein table on the third port holds almost the same information as the table in the **PIA Result Analysis** view. Node, that the accessions are lists again, as PIA always

reports protein groups. Additionally, each group has a clusterID, which represents the connected set or tree in the PIA intermediate structure. Elements in such a cluster are connected by their relations of accessions and PSMs, as can be seen in the **PIA Result Analysis** view on the top right. This does not mean, that reported protein groups with the same cluster ID share a peptide, but they are in the same component (a tree in the top right region of the **PIA Result Analysis** view) and can be connected by other (even not reported) accessions or peptides.

All of the reported tables can easily be processed with default KNIME nodes. This facilitates filtering, sorting, plotting etc. But also an analysis with R or Python can be created without much effort. If you like, you can also export the tables to excel and process them further.

3.2.5 The Export File Port

On this port the created export file can be found. This could be either used for storage (e.g. when creating an mzIdentML file) but also as input for OpenMS's **ProteinQuantifier**, if an idXML file was exported (as shown in the advanced tutorial).

3.3 Tasks

- Find the protein group "P02675" in the **PIA Results Analysis** view (it has a score of 46.66). Try to understand, why this protein (group) was reported, but no group with Q8K0E8 was reported.

- Now find find the protein group "Q91V41" (score 12.57). Why was this group and also the group "P62821, Q5SW88, Q9D1G1" (score 5.65) reported?

- Make a PIA analysis of the same data, but use an additional search engine. For this, you should activate the creation of PSM sets and set the used inference base score to the Combined FDR Score.
Hint: The workflow for this is given in the data folder, but try to extend the workflow on your own first. ;)

4 PIA for inference of quantitative data

PIA can be used for protein inference on quantitative data. For this, PIA can be allows to infere the protein groups and facilitate a quantification, which is not dependend on unique peptides. This is achieved due to the fact, that PIA groups the proteins, which have the same evidence, into a protein group. Be aware, though that protein-subgroups are not represented by this at the moment.

The inference for quantification works out of the box with PIA and OpenMS. A workflow showing this on a small dataset (just load in the provided files) is given in the data folder of this tutorial.

First you need to run a feature detection, map IDs, align and normalise the data to create a featureXML. More about these steps can be found in the OpenMS tutorial. Then, you can use PIA to create the inference input for the **PROTEINQUANTIFIER**. For this, you just need to select idXML and protein level as export in the general settings and run the **PIA ANALYSIS** node. The settings should be applied with respect to the provided data. All information for this can be found in the prior chapter.
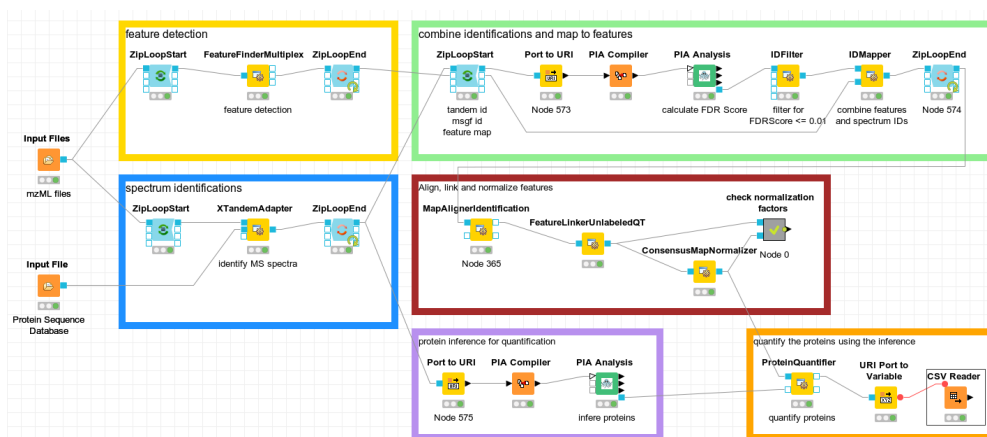
Figure 8: A basic workflow using PIA for the protein inference used for protein quantification.

4.1 Tasks

- Run the provided workflow on the small test dataset. Verify, that "P68871/contaminant_HBB_HUMAN" is a valid candidate for a regulated protein.

5 Web Frontend

The web frontend, which we did not discuss, can be tested at `https://github.com/BioDocker/containers/tree/master/pia-web-server/1.1.0-SNAPSHOT`.

References

[1] Uszkoreit et al., *PIA: An Intuitive Protein Inference Engine with a Web-Based User Interface.*, J Proteome Res, 2015.

[2] Jones et al., *Improving sensitivity in proteome studies by analysis of false discovery rates for multiple search engines.*, Proteomics, 2009.

[3] Perez-Riverol et al., *ms-data-core-api: an open-source, metadata-oriented library for computational proteomics.*, Bioinformatics, 2015.